

Appendix

Data Cleaning:

```
# Check for missing values in bicuphi
```

```
> summary(bicuphi)
```

```
      X      X.1      X.2
Length:1325 Length:1325 Length:1325
Class :character Class :character Class :character
Mode :character Mode :character Mode :character
```

```
> # More detailed check
```

```
> sapply(bicuphi, function(x) sum(is.na(x)))
```

```
  X X.1 X.2
0   0   0
```

~ Notes: There are no NA values in the data. That does not mean that there could be other types of data errors.

```
# Convert X.2 to numeric, coercing any non-numeric strings to NA
```

```
> bicuphi$X.2 <- as.numeric(as.character(bicuphi$X.2))
```

```
Warning message:
```

```
NAs introduced by coercion
```

```
# Check for newly created NAs in X.2 (if any non-numeric values were present)
```

```
> sum(is.na(bicuphi$X.2))
```

```
[1] 2
```

```
# Identify which rows have NAs in X.2 introduced by coercion
```

```
> na_rows <- which(is.na(bicuphi$X.2))
```

```
> print(na_rows)
```

```
[1] 1 2
```

```
# Optional: view the rows with NAs in X.2 to understand what the non-numeric values were
```

```
> print(bicuphi[na_rows, ])
```

```
      X  X.1 X.2
1              NA
2 DATE TIME  NA
```

```
# Correctly count empty strings in each column
```

```
> empty_string_counts <- sapply(bicuphi, function(x) sum(x == ""))
```

```
> print(empty_string_counts)
```

```
  X X.1 X.2
1   1   NA
```

```
      X  X.1 X.2
1              NA
2 DATE TIME  NA
```

```
> # Correcting the command to identify rows with empty strings
```

```
> empty_string_rows <- sapply(bicuphi, function(x) sum(x == ""))
```

```
> total_empty_strings <- sum(empty_string_rows)
```

```
> print(total_empty_strings)
```

```
[1] NA
```

```

>
> # Remove rows where X.2 is NA
> bicuphi_clean <- bicuphi[!is.na(bicuphi$X.2), ]
>
> # Check the dimensions to ensure rows are removed
> dim(bicuphi_clean)
[1] 1323    3

```

~ Notes: Two rows have been removed 1, 2.

```

# Renaming columns in the bicuphi_clean dataframe
> names(bicuphi_clean) <- c('date', 'time', 'demand')
>
> # Verify the changes
> head(bicuphi_clean)
      date time demand
3 1-Mar-05 6:30      1
4 1-Mar-05 6:45      2
5 1-Mar-05 7:00      4
6 1-Mar-05 7:15      0
7 1-Mar-05 7:30     10
8 1-Mar-05 7:45     13

```

~ Notes: Descriptive aspects of the columns have been resolved. Numeric type for demand is confirmed as numeric.

```

# Convert the 'date' column to Date format > bicuphi_clean$date <-
as.Date(bicuphi_clean$date, format="%d-%b-%y") > > # Assuming the 'time' column
is already in a suitable format, you might still want to combine 'date' and
'time' > # First, ensure 'time' is a character for concatenation >
bicuphi_clean$time <- as.character(bicuphi_clean$time) > > # Combine 'date' and
'time' into a single 'datetime' column > bicuphi_clean$datetime <-
as.POSIXct(paste(bicuphi_clean$date, bicuphi_clean$time), format="%Y-%m-%d
%H:%M") > > # Check if 'datetime' is correctly formatted and if time intervals
are consistent > summary(bicuphi_clean$datetime) Min. 1st Qu. Median Mean "2005-
03-01 06:30:00" "2005-03-06 10:22:30" "2005-03-11 14:15:00" "2005-03-11 14:15:00"
3rd Qu. Max. "2005-03-16 18:07:30" "2005-03-21 22:00:00"

```

~ Notes: Create a single datetime column for clarity and simplicity.


```

# Calculate differences between consecutive datetime entries
> time_diffs <- diff(bicuphi_clean$datetime)
>
> # Check if all differences are 15 minutes (900 seconds)
> all(time_diffs == 900) # Returns TRUE if all intervals are exactly 15 minutes
[1] FALSE


```

~ Notes: Not all intervals are 15 minutes

Data Visualization:



This is a time plot of the demand with a 'datetime' stamp on the x axis. Trend: Max usage is typically on Monday and decreases thru Sunday. Usage on weekdays builds steadily in the morning rush hour and peaks around the 5 pm rush hour.



Average daily usage shows 2 peak periods that revolve around the workday hours. There is a spread our period in the morning and more concentrated usage period around evening rush hour.

Naïve Forecast

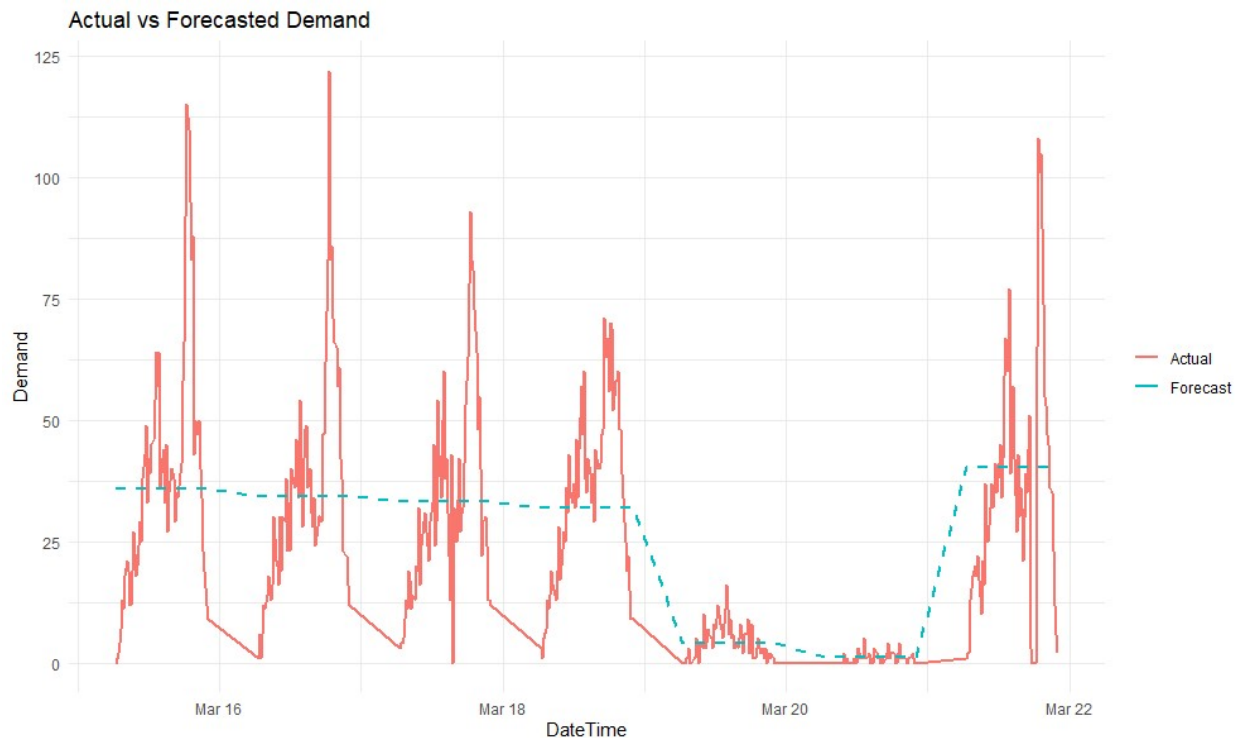
```
# Splitting the data
> training_set <- bicuphi_clean %>% filter(datetime < as.POSIXct("2005-03-15"))
> validation_set <- bicuphi_clean %>% filter(datetime >= as.POSIXct("2005-03-15"))
>
> # Calculating weekly patterns from the training data
> weekly_pattern <- training_set %>%
+   mutate(day_of_week = wday(datetime, label = TRUE)) %>%
+   group_by(day_of_week) %>%
+   summarise(average_demand = mean(demand))
>
> # Viewing the weekly pattern
> print(weekly_pattern)
# A tibble: 7 × 2
  day_of_week average_demand
  <ord>         <dbl>
1 Sun             1.12
2 Mon            40.3
3 Tue            35.9
4 Wed            34.2
5 Thu            33.2
6 Fri            31.9
7 Sat             4.11
>
> # Forecasting for the validation period
> validation_forecasts <- validation_set %>%
+   mutate(day_of_week = wday(datetime, label = TRUE)) %>%
+   left_join(weekly_pattern, by = "day_of_week") %>%
+   select(datetime, forecast_demand = average_demand)
>
> # Checking the forecasts
> head(validation_forecasts)
  datetime forecast_demand
1 2005-03-15 06:30:00      35.9127
2 2005-03-15 06:45:00      35.9127
3 2005-03-15 07:00:00      35.9127
4 2005-03-15 07:15:00      35.9127
5 2005-03-15 07:30:00      35.9127
6 2005-03-15 07:45:00      35.9127
>
> # Joining forecasts with actual demand
> validation_comparison <- validation_set %>%
+   select(datetime, actual_demand = demand) %>%
+   left_join(validation_forecasts, by = "datetime")
>
> # Calculating MAE
> mae <- mean(abs(validation_comparison$actual_demand -
validation_comparison$forecast_demand), na.rm = TRUE)
> print(paste("MAE:", mae))
[1] "MAE: 12.6568765072166"
>
> # Calculating RMSE
> rmse <- sqrt(mean((validation_comparison$actual_demand -
```

```

validation_comparison$forecast_demand)^2, na.rm = TRUE))
> print(paste("RMSE:", rmse))
[1] "RMSE: 18.9888449798315"
>
> # Ensure actual_demand and forecast_demand are numeric
> validation_comparison$actual_demand <-
as.numeric(validation_comparison$actual_demand)
> validation_comparison$forecast_demand <-
as.numeric(validation_comparison$forecast_demand)
>
> # Calculate MAPE
> mape <- mean(abs((validation_comparison$actual_demand -
validation_comparison$forecast_demand) / validation_comparison$actual_demand),
na.rm = TRUE) * 100
> print(paste("MAPE:", mape))
[1] "MAPE: Inf"

# Filtering out instances where actual_demand is zero to avoid division by zero
in MAPE calculation
> validation_comparison_filtered <- validation_comparison %>%
+   filter(actual_demand > 0)
>
> # Recalculating MAPE
> mape_filtered <- mean(abs((validation_comparison_filtered$actual_demand -
validation_comparison_filtered$forecast_demand) /
validation_comparison_filtered$actual_demand), na.rm = TRUE) * 100
> print(paste("Filtered MAPE:", mape_filtered))
[1] "Filtered MAPE: 127.594538853236"
>

```



Linear Regression Model:

```
# Summary of the model
> summary(lm_model)
```

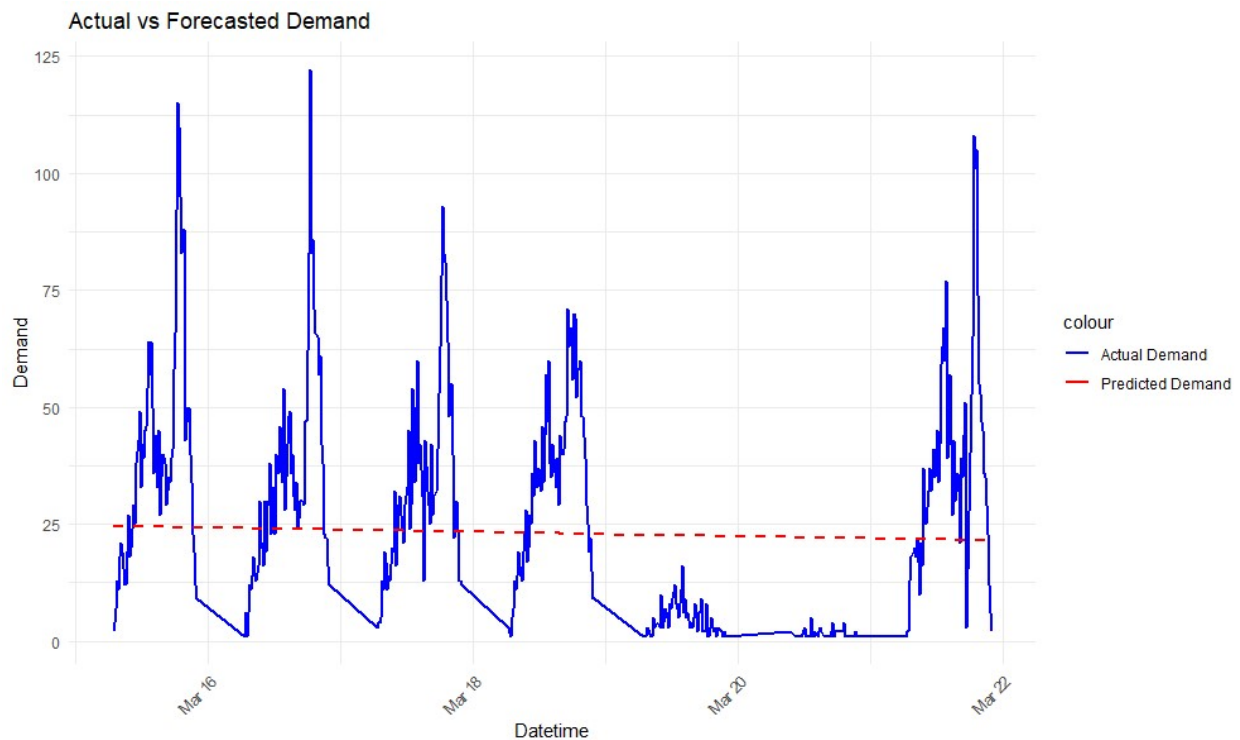
```
Call:
lm(formula = demand ~ datetime, data = training_set)
```

```
Residuals:
    Min     1Q   Median     3Q      Max
-29.989 -20.950  -1.847  13.339  107.994
```

```
Coefficients:
      Estimate Std. Error t value Pr(>|t|)
(Intercept)  5.917e+03  2.548e+03  2.323  0.0204 *
datetime    -5.305e-06  2.295e-06 -2.312  0.0210 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 22.9 on 814 degrees of freedom
Multiple R-squared:  0.006523,    Adjusted R-squared:  0.005302
F-statistic: 5.345 on 1 and 814 DF, p-value: 0.02104
```

```
# Print the error metrics
> print(paste("MAE:", mae))
[1] "MAE: 18.6840522177133"
> print(paste("MAPE:", mape))
[1] "MAPE: 314.343608815863"
> print(paste("RMSE:", rmse))
[1] "RMSE: 24.1018505149405"
```



Quadratic Model:

```
# Summary of the quadratic model  
> summary(lm_quadratic_model)
```

```
Call:  
lm(formula = demand ~ datetime_numeric + I(datetime_numeric^2),  
    data = training_set_filtered)
```

Residuals:

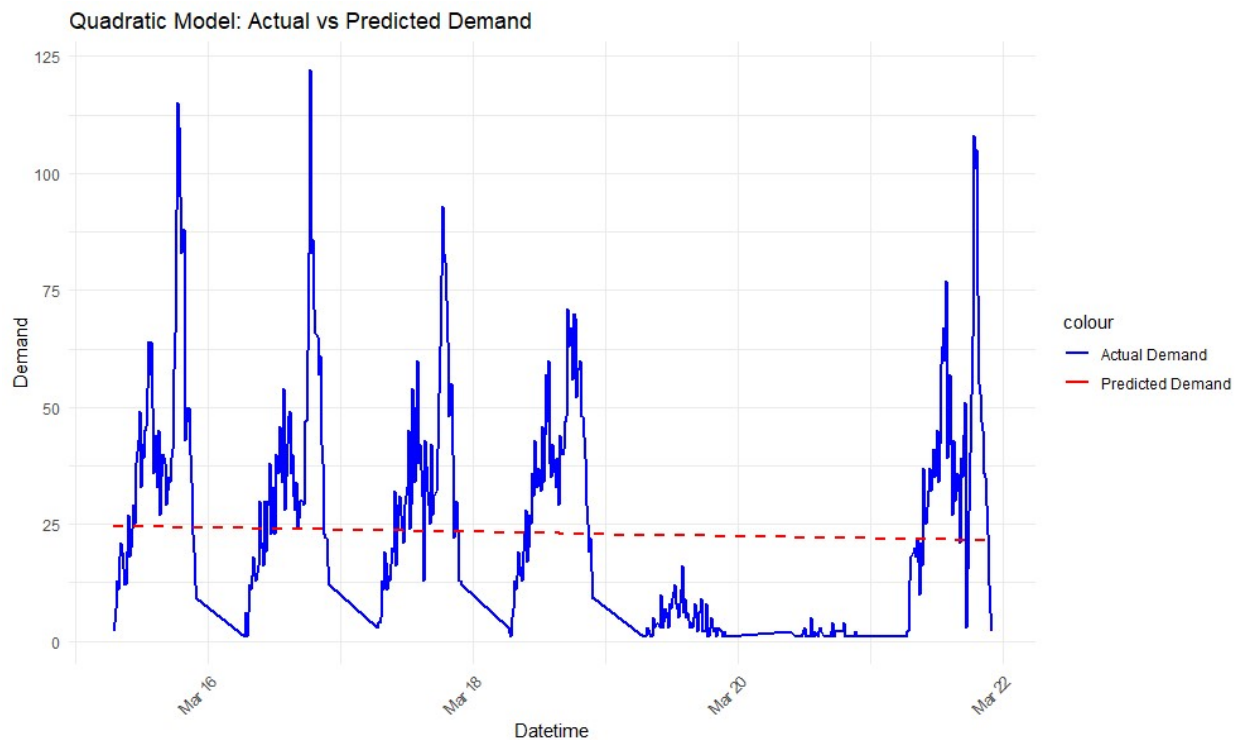
Min	1Q	Median	3Q	Max
-29.989	-20.950	-1.847	13.339	107.994

Coefficients: (1 not defined because of singularities)

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	5.917e+03	2.548e+03	2.323	0.0204 *
datetime_numeric	-5.305e-06	2.295e-06	-2.312	0.0210 *
I(datetime_numeric^2)	NA	NA	NA	NA

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 22.9 on 814 degrees of freedom
Multiple R-squared: 0.006523, Adjusted R-squared: 0.005302
F-statistic: 5.345 on 1 and 814 DF, p-value: 0.02104




Holt Winters Method:


```

> # Print the error metrics
> print(paste("MAE:", mae))
[1] "MAE: 39.7316740676646"
> print(paste("MAPE:", mape))
[1] "MAPE: 394.020267682534"
> print(paste("RMSE:", rmse))
[1] "RMSE: 49.0510276675861"
>


```



Holt-Winters Forecast vs Actual Demand



Holt-Winters Forecast vs Actual Demand



Arima

